

Data Cleaning & Preprocessing

Garbage In, Garbage Out — Fix It First

■ DATA QUALITY

■ LEARNING OUTCOME

By the end of this PDF, you will confidently identify and fix the most common data quality problems using Excel and Python, preparing any dataset for analysis.

■ Skills You Will Gain

- Identify missing, duplicate, and inconsistent data
- Handle missing values (drop, fill, impute)
- Remove duplicates and fix data types
- Normalize and standardize numeric data
- Encode categorical variables for analysis

SECTION 1

The Dirty Data Problem

Problem Type	Example	Impact
Missing Values	Age = NaN, Salary = blank	Wrong averages, model failures
Duplicates	Same order recorded twice	Inflated totals, false counts
Wrong Data Type	Date stored as text '01-Jan'	Cannot sort or calculate
Inconsistent Format	'Male', 'male', 'M' in Gender	Wrong group counts
Outliers	Salary = ■99,99,99,999	Skewed mean, bad predictions
Irrelevant Columns	CustomerID, Timestamp for name analysis	Noise in model

SECTION 2

Handling Missing Values

1	Detect df.isnull().sum() — count missing per column
2	Decide Is it <5%? Drop rows. Is it >20%? Drop column or impute.
3	Fill Numeric df['col'].fillna(df['col'].mean()) — fill with mean
4	Fill Categorical df['col'].fillna(df['col'].mode()[0]) — fill with most frequent
5	Advanced Use KNNImputer or median for skewed data

Python Code:

```
import pandas as pd
df = pd.read_csv('data.csv')
print(df.isnull().sum()) # Check missing
df['Age'].fillna(df['Age'].median(), inplace=True) # Fill with median
df.dropna(subset=['Salary'], inplace=True) # Drop rows with missing Salary
```

SECTION 3

Removing Duplicates

```
df.duplicated().sum() # How many duplicates?
df.drop_duplicates(inplace=True) # Remove all duplicates
df.drop_duplicates(subset=['OrderID'], keep='first') # By specific column
```

SECTION 4

Fixing Data Types

```
df.dtypes # Check all column types
df['Date'] = pd.to_datetime(df['Date']) # Convert to date
df['Price'] = df['Price'].astype(float) # Convert to number
df['Category'] = df['Category'].astype('category') # Memory efficient
```

SECTION 5

Standardization & Normalization

Technique	Formula	When to Use
Min-Max Normalization	$(x - \min) / (\max - \min) \rightarrow [0,1]$	When you need values between 0 and 1
Z-Score Standardization	$(x - \text{mean}) / \text{std} \rightarrow \text{mean}=0, \text{std}=1$	When data is normally distributed
Log Transform	$\log(x)$ — reduces skewness	Skewed data like salary, price

```
from sklearn.preprocessing import MinMaxScaler, StandardScaler
scaler = MinMaxScaler()
df[['Age', 'Salary']] = scaler.fit_transform(df[['Age', 'Salary']])
```

SECTION 6

Encoding Categorical Variables

```
# Label Encoding (for ordinal: Low, Medium, High)
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['Grade'] = le.fit_transform(df['Grade']) # One-Hot Encoding (for nominal: Red, Blue, Green)
df = pd.get_dummies(df, columns=['Color'], drop_first=True)
```

■ PRACTICE QUESTIONS

- Q1. What function in Python shows how many missing values are in each column?
Answer: df.isnull().sum()
- Q2. When should you use median instead of mean to fill missing values?
Answer: When the data is skewed (has outliers) — median is not affected by extreme values.
- Q3. What is the difference between normalization and standardization?
Answer: Normalization scales to [0,1]; standardization scales to mean=0, std=1.
- Q4. Why should you remove duplicate rows before analysis?
Answer: Duplicates inflate counts and totals, leading to incorrect analysis results.
- Q5. What is one-hot encoding and when do you use it?
Answer: Converting categorical variables into binary columns. Use for nominal (unordered) categories.

■ MINI PROJECT

Apply what you've learned with a real dataset!

Dataset: employee_dirty.csv (Name, Age, Salary, Department, Join_Date)

Steps:

1	Detect Run <code>df.isnull().sum()</code> and <code>df.duplicated().sum()</code>
2	Fix Missing Fill Age with median, Department with mode
3	Remove Dups Drop duplicate rows based on Name + Join_Date
4	Fix Types Convert Join_Date to datetime, Salary to float
5	Normalize Apply Min-Max scaling to Age and Salary
6	Encode One-hot encode Department column

Expected Output: A clean, analysis-ready DataFrame with no nulls or duplicates

Insights to Find: How many rows were removed? What % were dirty? Distribution of departments.

■ Pandas Data Cleaning Cheat Sheet	
<code>df.shape</code>	Rows and columns count
<code>df.info()</code>	Data types + null counts
<code>df.describe()</code>	Stats: mean, std, min, max
<code>df.isnull().sum()</code>	Missing values per column
<code>df.duplicated().sum()</code>	Count duplicate rows
<code>df.drop_duplicates()</code>	Remove duplicate rows
<code>df.fillna(value)</code>	Fill missing with value
<code>df.dropna()</code>	Remove rows with any null
<code>df.astype(type)</code>	Change column data type
<code>pd.to_datetime(col)</code>	Convert text to datetime
<code>df.rename(columns={})</code>	Rename columns
<code>df.drop(columns=[])</code>	Remove columns