

Hypothesis Testing

Make Data-Driven Decisions With Statistical Confidence

■ STATISTICAL TESTING

■ LEARNING OUTCOME

By the end of this PDF you will design, execute, and correctly interpret hypothesis tests to validate business decisions — from A/B tests to clinical trials to product analytics.

■ Skills You Will Gain

- Formulate null and alternative hypotheses correctly
- Choose the right statistical test for any situation
- Run t-tests, chi-square, ANOVA, and Mann-Whitney
- Understand Type I & II errors, power, and sample size
- Design and analyze A/B tests from scratch
- Report results in business language

■ Hypothesis testing is how Netflix decides which thumbnail to show, how Amazon chooses button colors, and how pharma companies approve drugs. This skill is everywhere.

SECTION 1

The Hypothesis Testing Framework

1	Define Hypotheses H_0 (null): No effect/difference exists H_A (alternative): An effect/difference exists
2	Choose α (significance level) $\alpha = 0.05$ (5% false positive rate) is industry standard Use $\alpha=0.01$ for medical/financial decisions
3	Choose the right test Based on: data type, number of groups, distribution shape
4	Calculate test statistic t-score, z-score, chi-square, F-statistic depending on test
5	Get p-value Probability of observing this result if H_0 is true
6	Make decision $p < \alpha \rightarrow$ Reject H_0 (result is significant) $p \geq \alpha \rightarrow$ Fail to reject H_0 (insufficient evidence)
7	Report finding Business language: 'Version B increased conversion by 12% ($p=0.02$)'

SECTION 2

Choosing the Right Test

Scenario	Test	Function
Compare mean to known value	One-sample t-test	<code>stats.ttest_1samp(data, popmean)</code>
Compare means: 2 independent groups	Two-sample t-test	<code>stats.ttest_ind(a, b)</code>
Compare means: 2 paired groups	Paired t-test	<code>stats.ttest_rel(before, after)</code>
Compare means: 3+ groups	One-way ANOVA	<code>stats.f_oneway(g1,g2,g3)</code>
Non-normal, 2 groups	Mann-Whitney U	<code>stats.mannwhitneyu(a,b)</code>
Non-normal, 3+ groups	Kruskal-Wallis	<code>stats.kruskal(g1,g2,g3)</code>
Two categorical variables	Chi-square test	<code>stats.chi2_contingency(table)</code>
Proportion comparison	Z-test for proportions	<code>proportions_ztest(count,nobs)</code>
Normality check	Shapiro-Wilk / K-S	<code>stats.shapiro(data)</code>

SECTION 3

Running Tests in Python

```
from scipy import stats import numpy as np # ■■ Two-sample t-test ■■ group_A =
df[df['Group']=='A']['Conversion'] group_B = df[df['Group']=='B']['Conversion'] t, p =
stats.ttest_ind(group_A, group_B, equal_var=False) # Welch's t-test print(f't={t:.3f}, p={p:.4f},
significant={p<0.05}') # ■■ Chi-square test ■■ ct = pd.crosstab(df['Gender'], df['Purchased'])
chi2, p, dof, expected = stats.chi2_contingency(ct) print(f'chi2={chi2:.3f}, p={p:.4f}') # ■■
One-way ANOVA ■■ groups = [df[df['Region']==r]['Sales'] for r in df['Region'].unique()] f, p =
stats.f_oneway(*groups) print(f'F={f:.3f}, p={p:.4f}') # ■■ Check normality first ■■ stat, p_norm
= stats.shapiro(group_A) if p_norm < 0.05: print('Non-normal → use Mann-Whitney') u, p_mw =
stats.mannwhitneyu(group_A, group_B, alternative='two-sided')
```

SECTION 4

Type I & II Errors + Statistical Power

	H True	H False
Reject H	Type I Error (α) — False Positive 'We said it works but it doesn't'	Correct Decision ■ (Power = $1-\beta$)
Fail to Reject H	Correct Decision ■	Type II Error (β) — False Negative 'We missed a real effect'

```
# Calculate required sample size for desired power from statsmodels.stats.power import
TTestIndPower analysis = TTestIndPower() n = analysis.solve_power(effect_size=0.5, # Cohen's d
power=0.80, # 80% chance of detecting real effect alpha=0.05) # 5% false positive rate
print(f'Need {int(np.ceil(n))} samples per group')
```

SECTION 5

A/B Testing End-to-End

1	Define Metric Primary: conversion rate. Secondary: revenue per user, session time
2	Calculate Sample Size Use power analysis: effect_size=0.05 (5% lift), power=0.8, α =0.05
3	Run Experiment Randomly split users 50/50, run for full business cycle (≥ 1 week)
4	Check for Bias Is randomization working? Compare baseline demographics of A vs B
5	Run t-test or z-test proportions_ztest([conv_A,conv_B],[n_A,n_B])
6	Check practical significance Statistical significance $\neq$ business significance. Is 0.1% lift worth the cost?

'Version B increased conversion from 3.2% to 3.8% (+18.75%), p=0.012 — implement'

```
# A/B test with proportions from statsmodels.stats.proportion import proportions_ztest conversions
= [120, 145] # B converted more nob = [2000, 2000] # same sample size z, p =
proportions_ztest(conversions, nob) print(f'z={z:.3f}, p={p:.4f}') lift = (145/2000 - 120/2000) /
(120/2000) * 100 print(f'Lift: {lift:.1f}%')
```

SECTION 6

Effect Size — Why p-value Alone Is Not Enough

Effect Size Measure	Formula	Interpretation
Cohen's d (t-test)	$d = (\mu_{\text{B}} - \mu_{\text{A}}) / \text{pooled_std}$	0.2=small, 0.5=medium, 0.8=large
Cramer's V (chi-square)	$V = \sqrt{(\chi^2 / n \cdot \min(r-1, c-1))}$	0.1=small, 0.3=medium, 0.5=large
η^2 (ANOVA)	$\eta^2 = \text{SS_between} / \text{SS_total}$	Proportion of variance explained

■ PRACTICE QUESTIONS

Q1. What does p-value = 0.03 mean?

Answer: If the null hypothesis were true, there's only a 3% chance of observing a result this extreme by chance. Since $0.03 < 0.05$, we reject H_0 .

Q2. You want to compare revenue across 4 sales regions. Which test do you use?

Answer: One-way ANOVA — it compares means across 3 or more independent groups simultaneously.

Q3. What is Type I error and how do you control it?

Answer: False positive — rejecting H_0 when it's actually true. Controlled by the significance level α . Setting $\alpha=0.05$ means you accept a 5% false positive rate.

Q4. Why is sample size calculation important before running an A/B test?

Answer: Too small = underpowered test (miss real effects). Too large = waste of time and money. Power analysis finds the minimum sample needed to detect a meaningful effect.

Q5. Statistical test shows p=0.04 for a 0.01% conversion lift. Should you ship the change?

Answer: No — it's statistically significant but practically insignificant. Always check effect size and business impact alongside p-value.

■ Hypothesis Testing Cheat Sheet	
H_0 : no effect	Null hypothesis — default assumption
H_1 : effect exists	Alternative hypothesis — what you test
$\alpha = 0.05$	5% significance level (standard)
$p < \alpha \rightarrow \text{reject } H_0$	Result is statistically significant
<code>stats.ttest_ind(a,b)</code>	2-group mean comparison
<code>stats.ttest_rel(b,a)</code>	Paired before/after test
<code>stats.f_oneway(*groups)</code>	ANOVA: 3+ groups
<code>stats.chi2_contingency(t)</code>	Chi-square: categorical
<code>stats.mannwhitneyu(a,b)</code>	Non-parametric 2-group

<code>stats.shapiro(data)</code>	Normality check
<code>proportions_ztest(c,n)</code>	A/B test proportions
Cohen's d = 0.8	Large effect size